

# Working with Data

July 6, 2019

## 1 Working with Data

The purpose of this chapter is to familiarize the reader with some of the basic of working with data in Julia. As would be expected, much of the focus of this chapter is on or around dataframes, including dataframe functions. Other topic covered include categorical data, input-output (IO), and the split-apply-combine strategy.

### 1.1 Dataframes

A dataframe is a tabular representation of data, similar to a spreadsheet or a data matrix. As with a matrix, the observations are rows and the variables are columns. Each row is a single (vectorvalued) observation. For a single row, i.e., observation, each column represents a single realization of a variable. At this stage, it may be helpful to explicitly draw the analogy between a dataframe and the more formal notation often used in statistics and data science.

Suppose we observe  $n$  realizations  $\mathbf{x}_1, \dots, \mathbf{x}_n$  of  $p$ -dimensional random variables  $\mathbf{X}_1, \dots, \mathbf{X}_n$ , where  $\mathbf{X}_i = (X_{i1}, X_{i2}, \dots, X_{ip})'$  for  $i = 1, \dots, n$ . In matrix form, this can be written

$$\mathcal{X} = (\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_n)' = \begin{pmatrix} \mathbf{X}'_1 \\ \mathbf{X}'_2 \\ \vdots \\ \mathbf{X}'_n \end{pmatrix} = \begin{pmatrix} X_{11} & X_{12} & \cdots & X_{1p} \\ X_{21} & X_{22} & \cdots & X_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ X_{n1} & X_{n2} & \cdots & X_{np} \end{pmatrix}. \quad (1)$$

Now,  $\mathbf{X}_i$  is called a random vector and  $\mathcal{X}$  is called an  $n \times p$  random matrix. A realization of  $\mathcal{X}$  can be considered a data matrix. For completeness, note that a matrix  $\mathbf{A}$  with all entries constant is called a constant matrix.

Consider, for example, data on the weight and height of 500 people. Let  $\mathbf{x}_i = (x_{i1}, x_{i2})'$  be the associated observation from the  $i$ th person,  $i = 1, 2, \dots, 500$ , where  $x_{i1}$  represents their weight and  $x_{i2}$  represents their height. The associated data matrix is then

$$\mathcal{X} = (\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_{500})' = \begin{pmatrix} x'_1 \\ x'_2 \\ \vdots \\ x'_{500} \end{pmatrix} = \begin{pmatrix} x_{11} & x_{12} \\ x_{21} & x_{22} \\ \vdots & \vdots \\ x_{500,1} & x_{500,2} \end{pmatrix} \quad (2)$$

A dataframe is a computer representation of a data matrix. In Julia, the `DataFrame` type is available through the `DataFrames.jl` package. There are several convenient features of a `DataFrame`, including:

- columns can be different Julia types;
- table cell entries can be missing;
- metadata can be associated with a DataFrame;
- columns can be names;
- tables can be subsetted by row, column or both.

The columns of a DataFrame are most often integers, floats or strings, and they are specified by Julia symbols.

```
In [1]: ## symbol versus string
        fruit = "apple"
        println("eval(:fruit): ", eval(:fruit))
```

```
eval(:fruit): apple
```

```
In [2]: println("""eval("apple"): """, eval("apple"))
```

```
eval("apple"): apple
```

In Julia, a symbol is how a variable name is represented as data; on the other hand, a string represents itself. Note that `df[:symbol]` is how a column is accessed with a symbol; specifically, the data in the column represented by `symbol` contained in the DataFrame `df` is being accessed. In Julia, a DataFrame can be built all once in multiple phases.

```
In [6]: ## Some examples with DataFrames
```

```
using DataFrames, Distributions, StatsBase, Random
```

```
Random.seed!(825)
```

```
N=50
```

```
## Create a sample dataframe
```

```
## Initially the DataFrame has N rows and 3 columns
```

```
df1 = DataFrame(
    x1 = rand(Normal(2,1),N),
    x2 = [sample(["High", "Medium", "Low"],
                pweights([0.25,0.45,0.30])) for i = 1:N],
    x3 = rand(Pareto(2,1),N)
)
```

```
## Add a 4th column, y, which is dependent on x3 and the level of x2
```

```
df1[:y] = [df1[i,:x2] == "Medium" ? *(2, df1[i, :x3]) :
            df1[i,:x2] == "High" ? *(4, df1[i, :x3]) :
            *(0.5, df1[i,:x3]) for i=1:N]
```

```
Out[6]: 50-element Array{Float64,1}:
```

```
2.393325358524223
2.4310412249328515
5.0422956234853125
3.7345079211580465
0.8588031634040387
0.5231696808320803
2.031479078828272
2.4806651435260076
9.576090582790384
6.124545507841852
0.7720685108626731
6.192654315744594
6.699890875731672

3.691184458092902
9.562195756696816
2.433614543200589
2.412383239462625
1.222717094180467
3.6746066646508866
2.5425484862899155
0.6933323954399253
7.280430172887931
3.2074781385911297
4.770360925053668
2.39414636559205
```

A DataFrame can be sliced the same a two-dimensional Array is sliced, i.e., `df[row_range, column_range]`. These ranges can be specified in a number of ways:

- Using Int indices individually or as arrays, e.g., 1 or [4,6,9].
- Using `:` to select indicis in a dimension, e.g., `x:y` selects the range from x to y and `:` selects all indices in that dimension.
- Via arrays of Boolean values, where `true` selects the elements at that index.

Note that columns can be selected by their symbols, either individually or in an array `[:x1, :x2]`.

```
In [7]: ## Slicing DataFrames
println("df1[1:2,3:4]: ",df1[1:2, 3:4])
```

```
df1[1:2,3:4]: 2x2 DataFrame
```

```
Row  x3      y
     Float64  Float64
```

```
1    1.19666  2.39333
2    1.21552  2.43104
```

```
In [8]: println("\ndf1[1:2, [:y,:x1]]: ",df1[1:2, [:y,:x1]])
```

```
df1[1:2, [:y,:x1]]: 2E2 DataFrame
```

```
Row y          x1
    Float64  Float64
```

```
1    2.39333  1.81025
2    2.43104  3.1707
```

```
In [9]: ## Now, exclude columns x1 and x2
        keep = setdiff(names(df1), [:x1, :x2])
        println("\nColumns to keep", keep)
```

```
Columns to keepSymbol[:x3, :y]
```

```
In [10]: println("df1[1:2,keep]: ",df1[1:2, keep])
```

```
df1[1:2,keep]: 2E2 DataFrame
```

```
Row x3          y
    Float64  Float64
```

```
1    1.19666  2.39333
2    1.21552  2.43104
```

In practical applications, missing data is common. In `DataFrames.jl`, the `Missing` type is used to represent missing values. In Julia, a singleton occurrence of `Missing`, `missing` is used to represent missing data. Specifically, `missing` is used to represent the value of a measurement when a valid value could have been observed but was not. Note that `missing` in Julia is analogous to `NA` in *R*.

In the following code block, the array `v2` has type `Union{Float64, Missings.Missing}`. In Julia, `Union` types are an abstract type that contain objects of types included in its arguments. In this example, `v2` can contain values of `missing` or `Float64` numbers. Note that `missings()` can be used to generate arrays that will support missing values; specifically, it will generate vectors of type `Union` if another type is specified in the first argument of the function call. Also, `ismissing(x)` is used to test whether `x` is missing, where `x` is usually an element of a data structure, e.g., `ismissing(v2[1])`.

```
In [13]: v1 = missings(2)
        println("v1: ",v1)
```

```
v1: Missing[missing, missing]
```

```
In [14]: v2 = missings(Float64,1,3)
        v2[2] = pi
        println("v2: $(v2)")
```

```
v2: Union{Missing, Float64}[missing 3.14159 missing]
```

```
In [15]: ## test for missing  
         m1 = map(ismissing,v2)  
         println("m1: $(m1)")
```

```
m1: Bool[true false true]
```

```
In [16]: println("Percent missing v2: ", *(mean([ismissing(i) for i in v2]), 100))
```

```
Percent missing v2: 66.66666666666666
```

Note that most functions in Julia do not accept data of type `Missings.Missing` as input. Therefore, users are often required to remove them before they can use specific functions. Using `skipmissing()` returns an iterator that excludes the missing values and, when used in conjunction with `collect()`, gives an array of non-missing values. This approach can be used with functions that take non-missing values only.

```
In [17]: ## calculates the mean of the non-missing values  
         mean(skipmissing(v2))
```

```
Out[17]: 3.141592653589793
```

```
In [18]: ## collects the non-missing values in an array  
         collect(skipmissing(v2))
```

```
Out[18]: 1-element Array{Float64,1}:  
          3.141592653589793
```

## 1.2 Categorical Data

In Julia, categorical data is represented by arrays of type `CategoricalArray`, defined in the `CategoricalArrays.jl` package. Note that `CategoricalArray` arrays are analogous to factors in R. `CategoricalArray` arrays have a number of advantages over `String` arrays in a dataframe:

- They save memory by representing each unique value of the string array as an index.
- Each index corresponds to a level.
- After data cleaning, there are usually only a small number of levels.

`CategoricalArray` arrays support missing values. The type `CategoricalArray{Union{T, Missing}}` is used to represent missing values. When indexing/slicing arrays of this type, `missing` is returned when it is present at that location.

```
In [22]: ## Number of entries for the categorical Arrays  
         Nca = 10  
  
         ## Empty array
```

```

v3 = Array{Union{String,Missing}}(undef, Nca)

##Array has string and missing values
v3 = [isodd(i) ? sample(["High", "Low"], pweights([0.35,0.65])) :
      missing for i = 1:Nca]
## v3c is of type CategoricalArray{Union{Missing,String},1,UInt32}
v3c = categorical(v3)

## Levels should be ["High","Low"]
println("1. levels(v3c): ", levels(v3c))
println("1. v3c: $(v3c)")

1. levels(v3c): ["High", "Low"]
1. v3c: Union{Missing, CategoricalString{UInt32}}["Low", missing, "High", missing, "Low", missi

In [20]: ## Reordered levels - does not change the data
         levels!(v3c, ["Low", "High"])
         println("2.levels(v3c): $(levels(v3c))")

2.levels(v3c): ["Low", "High"]

In [21]: println("2.v3c: $(v3c)")

2.v3c: Union{Missing, CategoricalString{UInt32}}["High", missing, "Low", missing, "Low", missi

```

Here are several useful functions that can be used with CategoricalArray arrays:

- `levels()` returns the levels of the CategoricalArray.
- `levels()` changes the order of the array's levels.
- `compress()` compresses the array saving memory.
- `decompress()` decompresses the compressed array.
- `categorical(ca)` converts the array `ca` into an array of type CategoricalArray.
- `droplevels!(ca)` drops levels no longer present in the array `ca`. This is useful when a dataframe has been subsetting and some levels are no longer present in the data.
- `recode(a, pairs)` recodes the levels of the array. New levels should be the same type as the original ones.
- `recode!(new, orig, pairs)` recodes the levels in `orig` using the `pairs` and puts the new levels in `new`.